

From dlt traces to a **1.7B** model that replaces your **744B** agent

Four steps, no pipeline to build. Same behaviour on the task, but 50-90% lower inference cost.

Point dlt at your traces

A dltHub verified source pulls your agent traces from where they already live, Pydantic, Arize, or Langfuse, and normalizes every vendor format to one OTel standard.

See what each agent costs

Agent Cost & Usage breaks AI spend down per model, per person, per customer. Pick one high-volume, well-scoped agent that is expensive to run. That is your candidate.

Turn its traces into a model

The **agent_distillation** transform lands a training-ready dataset. distil labs trains and hosts the model on validated synthetic data drawn from your traces, not the raw traces, which is why a small model can match the big one.

Swap one URL

The model is served behind the same OpenAI-compatible API the 744B agent used. You change one endpoint. Same behaviour on the task, 50-90% cheaper, and no training script, pipeline, or inference to operate.

It holds, even on degraded traces

