

The most **tunable** model isn't the best out of the box.


A mid-size MoE benchmark: four models, seven tasks. Fine-tuned, NVIDIA's Nemotron-3-Nano tied for the best.

A model's out-of-the-box score measures what it already knows, not how much it can learn. Those are different things, and only one of them survives fine-tuning.

So we measured how much each model can learn.

Base rank **doesn't** **predict fine-tuned** **rank.**

Average rank across seven tasks,
before and after fine-tuning. Lower is
better.



Fine-tuned, NVIDIA's Nemotron-3-Nano-30B matched its teacher on most of the seven tasks, running on just 3B active parameters.

18x

smaller than that teacher, Nemotron-3-Ultra-550B, and far cheaper to serve.

Tunability is as real a metric as any benchmark.

Models get scored out of the box. We think labs should report how well they fine-tune, too. That is the number that matters when you build on them.